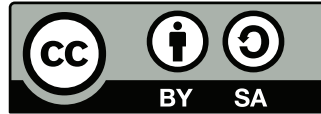


Operating Systems and Middleware: Supporting Controlled Interaction

Max Hailperin
Gustavus Adolphus College

Revised Edition 1.1
July 27, 2011

Copyright © 2011 by Max Hailperin.



This work is licensed under the Creative Commons Attribution-ShareAlike 3.0 Unported License. To view a copy of this license, visit

<http://creativecommons.org/licenses/by-sa/3.0/>

or send a letter to Creative Commons, 171 Second Street, Suite 300, San Francisco, California, 94105, USA.

Chapter 10

Messaging, RPC, and Web Services

10.1 Introduction

Application programmers who create distributed systems of communicating processes fundamentally rely upon the support for networking provided by operating systems; this support was described in Chapter 9. Sometimes this reliance is direct; for example, the application programs may use sockets to open connections between TCP ports and then send byte streams that encode application messages. Increasingly, however, this reliance is indirect because a middleware layer comes between the application program and the socket API. The application programmer works in terms of the middleware-supported abstractions, such as message queues and remotely accessible objects. The middleware provides those abstractions by making use of the more fundamental operating system-supported facilities.

In this chapter, I will present two distinct styles of communication middleware. Messaging systems, discussed in Section 10.2, support the one-way transmission of messages. Application programmers can choose to use those messages in request-response pairs, but the middleware remains oblivious to this pairing and treats each message individually. Sending a request and receiving a response happen in separate transactions, potentially at quite different times. For more tightly coupled interactions, Remote Procedure Call (RPC) systems provide an alternative style of communication, as presented in Section 10.3. Each time a process invokes an RPC operation, it sends a request and then immediately waits for a response as part of the same transaction. The RPC system makes the request-response pair appear

to the application program as a normal procedure call and return.

After presenting each of these styles of communication, I turn in Section 10.4 to their connection with web services. Web services use standardized communication mechanisms to make programmed functionality available over the Internet. Most web services fit the RPC model, though they can also use one-way messaging or (in theory) more general message exchange patterns.

Finally, the chapter concludes, as usual, with a look at security issues in Section 10.5 and then exercises, projects, and notes.

10.2 Messaging Systems

Applications based on one-way transmission of messages use a form of middleware known as *messaging systems* or *message-oriented middleware (MOM)*. One popular example of a messaging system is IBM's WebSphere MQ, formerly known as MQSeries. One popular vendor-neutral API for messaging is the Java Message Service (JMS), which is part of J2EE.

Messaging systems support two different forms of messaging: *message queuing* and *publish/subscribe messaging*. I already introduced message queuing in Section 5.2.2. Here I will build on that introduction to message queuing and also provide an introduction to publish/subscribe messaging.

Figure 10.1 illustrates the difference between the two forms of messaging. Message queuing strongly decouples the timing of the client and the server, because the queue will retain the messages until they are retrieved. (Optionally, the client can specify an expiration time for unretrieved messages.) The server need not be running at the time a message is sent. On the other hand, the client is only weakly decoupled from the server's identity. Although the client doesn't send the message to a specific server, it does send it to a specific queue, which still creates a *point-to-point* architectural structure, because each queue normally functions as the in-box for a particular server. A point-to-point structure means that if the message is of interest to multiple servers, the client needs to send it to multiple queues. The publish/subscribe architecture, in contrast, strongly decouples publishers from any knowledge of the subscribers' identities. Each message is sent to a general topic and from there is distributed to any number of subscribers that have indicated their interest in the topic. However, publish/subscribe messaging usually does not strongly decouple timing. Messages are usually only sent to current subscribers, not retained for future subscribers.

The portion of a messaging system managing topic objects for pub-




Figure 10.1: Message queuing involves three steps: (1) the client sends a message to a queue, (2) the queue retains the message as long as necessary, (3) the server retrieves the message. Publish/subscribe messaging involves a different sequence of three steps: (1) each subscriber subscribes with one of the messaging system’s “topic” objects, (2) the publisher sends a message to the topic, (3) the message is distributed to all current subscribers.

lish/subscribe messaging is known as a *broker*. The broker is responsible for maintaining a list of current subscribers for each topic and for distributing each incoming publication to the current subscribers of the publication’s topic.

Section 5.2.2 explained the relationship between message queuing and transactions. A transaction can retrieve messages from queues, do some processing, such as updating a database, and send messages to queues. When the transaction commits, the input messages are gone from their queues, the database is updated, and the output messages are in their queues. If the transaction aborts, then the input messages remain in their queues, the database remains unchanged, and the output messages have not entered their queues.

This transactional nature of message queuing has an important consequence for systems in which request messages are paired with response messages and will help me explain the difference between messaging and RPC systems. Consider a client and server coupled through request and response queues, as shown in Figure 10.2. The client can generate a request message in one transaction and then in a second transaction wait for a response message. However, it cannot do both in a single transaction, or it will wait forever, having deadlocked itself. Until the transaction commits, the request message doesn’t enter the request queue. As a result, the server has nothing to respond to and won’t generate a response message. Therefore, the client will continue waiting for the response message and so the transaction won’t




Figure 10.2: A client and server can engage in a request-response protocol using two message queues. Typically, the client tags each request message with a unique identifying string, known as a *correlation ID*. The server copies this ID into the resulting response message so that the client knows to which request the response corresponds.

commit, completing the deadlock. If your goal is to have the client make use of the server as one indivisible step within a transaction, then you need to use RPC rather than messaging.

Publish/subscribe messaging can participate in transactions as well, but the results are less interesting. Publishing is just like sending to a queue, in that the message isn't actually sent until the transaction commits. However, receipt of messages by subscribers is handled differently. If a subscriber receives a message within a transaction and then aborts the transaction, it cannot count on the message being redelivered when the transaction is retried.

In either messaging model, a consumer may want to receive only selected messages that are of interest to it. For example, it may want to receive stock ticker messages with updated prices, but only for IBM stock and only if the price is less than 75 or more than 150. The program could receive all stock ticker messages (by reading from a queue to which they are sent or by subscribing to a topic to which they are published) and ignore those that are uninteresting. However, for the sake of efficiency, messaging systems generally provide mechanisms to do the filtering prior to message delivery.

In the publish/subscribe model, the selection of just IBM stock might be accomplished simply by having a sufficiently specific topic. Messaging systems generally allow topics to be hierarchically structured, much like files in directory trees or Internet domains within the DNS. Thus, a topic for IBM stock prices might be `finance/stockTicker/IBM`. A subscriber interested only in this one stock could subscribe to that specific

topic, whereas a subscriber interested in all stock prices could subscribe to `finance/stockTicker/+`, where the wildcard `+` indicates any one subtopic. Another wildcard, `#`, is fancier than needed in this case but can be useful in other circumstances. A subscription to `finance/stockTicker/#` would receive not only messages about each individual stock, such as IBM, but also general messages, directed to `finance/stockTicker` itself, and more specialized messages, directed to descendant subtopics any number of levels below `finance/stockTicker/IBM` and its siblings.

This hierarchy of topics is limited, however. It fits the publish/subscribe model but not the message queuing model, and it addresses only qualitative selection criteria that naturally lead to distinct topics. In the example I gave earlier, it is unlikely that a system architect would create three subtopics of IBM for `under75`, `between75and150`, and `over150`. Among other reasons, there may be other subscribers interested in other price ranges.

Therefore, messaging systems allow message consumers to specify more general selection criteria. In the JMS API, for example, if `s` is a messaging `Session` and `d` is a messaging `Destination`, that is, either a `Queue` or a `Topic`, then executing

```
s.createConsumer(d, "Symbol = 'IBM' AND " +
                    "(Price < 75 OR Price > 150)")
```

will produce a `Consumer` object with the specified selector. Any `receive` operation performed on that `Consumer` (or any `MessageListener` registered with that `Consumer`) will see only those messages satisfying the selection criterion.

10.3 Remote Procedure Call

The goal of *Remote Procedure Call* (*RPC*) middleware is to make request-response communication as straightforward for application programmers to program as ordinary procedure calls. The client application code calls a procedure in the ordinary way, that is, passing in some arguments and obtaining a return value for its further use. The procedure it calls just happens to reside in a separate server. Behind the scenes, the middleware encodes the procedure arguments into a request message and extracts the return value from the response message. Similarly, the server application code can take the form of an ordinary procedure. By the time it is invoked, the procedure arguments have already been extracted from the request message, freeing it from that responsibility. Section 10.3.1 explains further the principles

upon which RPC operates. Section 10.3.2 provides a concrete example of using RPC in the particular form of Java RMI (Remote Method Invocation). The subsequent section, 10.4, is devoted to web services but also provides additional information on how RPC plays out in that context.

10.3.1 Principles of Operation for RPC

To understand how RPC middleware functions, it is helpful to think about the fact that different procedures can present the same interface. For example, consider procedures for squaring a number. You could have several different procedures that take a numeric argument, compute the square, and return it. One might work by multiplying the number by itself. Another might use a fancy calculation involving logarithms. And a third might open up a network instant-messaging connection to a bored teenager, ask the teenager what the square of the number is, then return the value it receives, correctly extracted from the textual instant-messaging response. This third procedure is known as a *proxy* for the teenager. The proxy's method of squaring the number involves neither multiplication nor logarithms, but rather delegation of responsibility. However, the proxy still is presenting the same interface as either of the other two procedures.

Figure 10.3 shows how RPC middleware uses a proxy to put the client in the position of making a normal procedure call. The client application code actually does make a normal procedure call; that much is no illusion. However, it only gives the illusion of calling the server procedure that does the real computation. Instead, the called procedure is a proxy standing in for the server procedure; the proxy is often known as a *stub*. The stub proxy discharges its responsibility not by doing any actual computation itself, but by using request and response messages to communicate with the server.

The stub proxy suffices to hide communication issues from the application programmer writing the client code. In some cases, that is all that is needed, and the server is written by a networking expert who can directly write code to handle request and response messages. More typically, however, the server code is written by another application programmer who appreciates middleware support. As shown in Figure 10.4, the server application code can be a normal procedure, called the same way it would be if it were running on the same machine with the client. Once again, the illusion is only partial. The server application code really is being called with an ordinary procedure call. The only illusion concerns what code is doing the calling. From an application standpoint, the caller seems to be the client. However, the caller really is a dedicated portion of the RPC




Figure 10.3: In Remote Procedure Call, application code makes a normal procedure call to a stub proxy, which doesn't carry out the requested computation itself, but rather sends a request message to the server and then turns the response message into the procedure return value.




Figure 10.4: In order for the server application code to be free from communication details, it can be a normal procedure invoked by a portion of the RPC runtime sometimes called a skeleton or a tie.

runtime system, known as a *skeleton* or a *tie*, the purpose of which is to call the procedure in response to the request message. See Figure 10.5 for the application programmer's view of the result; the middleware communication disappears from view and the client application code seems to be directly calling the server application procedure, as though they were part of a single system.

Early versions of the RPC communication model were based on ordinary procedure calls, whereas more recent versions are based on the object-oriented concept of method invocation. The basic principles are the same, however, and the name RPC is commonly understood to include method invocation as well as procedure calling.

A key example of a non-object-oriented RPC standard is Open Network Computing (ONC) RPC, which was developed at Sun Microsystems and became an Internet standard. ONC RPC serves as the foundation for NFS, the Network File System discussed in Section 9.2.3. Each NFS operation, such as reading from a file, is carried out by calling an RPC stub procedure, which takes responsibility for packaging the procedure arguments into a request message and for extracting the return value from the response message.

In object-oriented versions of RPC, the stub proxy is an object presenting the same interface as the server object. That is, the same repertoire of methods can be invoked on it. The stub uses a uniform strategy for handling all methods: it translates method invocations into appropriate request messages.

Two significant object-oriented RPC standards are *CORBA* and *RMI*. *CORBA* (*Common Object Request Broker Architecture*) is a complicated language-neutral standard that allows code written in one language to call code written in another language and located elsewhere in a distributed sys-




Figure 10.5: The application programmer's view of an RPC system has the client code apparently making a direct call to the server procedure; the RPC proxy mechanism is invisible.

tem. RMI (*Remote Method Invocation*) is a considerably simpler mechanism included as part of the Java standard API; part of its simplicity comes from needing to support only a single programming language. Because RMI can optionally use CORBA's communication protocol, the Internet Inter-ORB Protocol (IIOP), the two systems can interoperate.

One important feature of object-oriented RPC systems such as RMI is that the values communicated as method arguments and return values can include references to other objects. That is, a remotely invoked method can operate not only on basic values, such as integers or strings, but also on user-defined types of objects. One remotely accessible object can be passed a reference that refers to another remotely accessible object. In fact, this is the way most objects find out about other objects to communicate with after getting past an initial startup phase.

To initially get communication started, client objects typically look up server objects using a *registry*, which is a specialized server that maintains a correspondence between textual names and references to the remotely accessible objects. (These correspondences are known as *bindings*.) The registry itself can be located, because it listens for connections on a prearranged port. When an application server object is created that a client might want to locate, the server object is bound to a name in the registry. The client then presents the same textual name to the registry in a lookup operation and thereby receives a reference to the initial server object it should contact.

After this initial contact is made, the objects can use the arguments and return values of remote method invocations to start passing each other references to additional objects that are not listed in the registry. Pretty soon any number of client and server objects can have references to one another and be invoking methods on each other.

Section 10.3.2 is occupied by an RMI programming example designed to reinforce the aforementioned point, that remote objects are located not only using the registry, but also by being passed as references. This example illustrates the way in which an application programmer uses RMI and thereby complements the preceding general discussion of how RPC stubs and skeletons work. I do not provide any more detailed information on the inner workings of RMI. However, in Section 10.4, I show the how RPC messages are formatted in the web services environment.

10.3.2 An Example Using Java RMI

Using RMI, it is possible to develop an implementation of the publish/subscribe messaging model, in which publishers send messages to topic objects, which

forward the messages along to subscribers. The code in this section shows such an implementation in the simplest possible form. In particular, this code has the following limitations; to address each limitation, there is at least one corresponding Programming Project:

- The message delivery is fully synchronous. That is, the publisher asks the topic object to deliver a message; control does not return to the publisher until the message has been delivered to all the subscribers. Programming Projects 10.2 and 10.3 address this.
- The example programs support only a single topic. Programming Projects 10.4 and 10.5 address this.
- In the example code, there is no way for a subscriber to explicitly unsubscribe from a topic. However, the code does support subscribers that terminate, lose communication, or otherwise fail. Programming Project 10.6 provides explicit unsubscription.
- The example code includes simple command-line interfaces for sending textual strings as messages and for displaying received messages on the terminal. These suffice to demonstrate the communication but do not have the appeal of a chat-room application or multi-player game. Programming Project 10.7 provides the opportunity to address this shortcoming.

When using RMI, each object that is remotely accessible must implement a Java interface that extends `java.rmi.Remote`. Each method in that interface must be declared as potentially throwing `java.rmi.RemoteException`. This potential for an exception is necessary because even if the underlying operation cannot possibly fail, the remote invocation of that operation can fail in myriad ways, such as through a network disconnection or a crash of the machine on which the remote object is located. Figure 10.6 shows the source code for a simple remote interface implemented by subscribers and also by topic objects. The reason why these two categories of participants in the publish/subscribe model implement this same interface is that they have something in common: they both receive messages.

Subscribers directly implement the `MessageRecipient` interface, as you will see later. However, topic objects need to implement an extension of the interface, because they can do more than receive messages; they can also add subscribers. Figure 10.7 shows the `Topic` interface, which extends `MessageRecipient` through the addition of a `addSubscriber` method. Notice that the argument passed to `addSubscriber` is itself a `MessageRecipient`.

```
import java.rmi.Remote;
import java.rmi.RemoteException;

public interface MessageRecipient extends Remote {
    void receive(String message) throws RemoteException;
}
```

Figure 10.6: The `MessageRecipient` interface describes the common feature shared by subscribers and the central topic objects that redistribute published messages to subscribers: any of these objects can receive a message.

This allows a reference to one remotely accessible object (the subscriber) to be passed to another (the topic) for its later use.

Having seen these two interfaces for remotely accessible objects, you are now ready to see an example of code that makes use of such an object. Figure 10.8 contains a simple program for sending a textual message (given as the first command-line argument) to a remote object implementing the `Topic` interface. The specific remote object is looked up with the aid of a registry, that is, a service within RMI that records name/object bindings. The registry is located on a server computer whose hostname is specified as the second command-line argument or on the local computer if no hostname is given.

Let's turn next to an example of how a remotely accessible object can be created and listed in the registry. The `Topic Server` class, as shown in Figures 10.9 and 10.10, implements the `Topic` interface. Each `TopicServer` keeps track of its current subscribers; additions happen in the `addSubscriber`

```
import java.rmi.RemoteException;

public interface Topic extends MessageRecipient {
    void addSubscriber(MessageRecipient subscriber)
        throws RemoteException;
}
```

Figure 10.7: The `Topic` interface provides an operation for subscribers to use to register their interest in receiving messages. By extending the `MessageRecipient` interface, the `Topic` interface is also prepared to receive messages from publishers.

```
import java.rmi.registry.LocateRegistry;
import java.rmi.registry.Registry;

public class Publisher {

    public static void main(String[] args) {
        if(args.length < 1 || args.length > 2){
            System.err.println("Usage: java Publisher message [host]");
            System.exit(1);
        }
        String message = args[0];
        String hostname = args.length > 1 ? args[1] : null;
        try {
            Registry registry = LocateRegistry.getRegistry(hostname);
            Topic topic = (Topic) registry.lookup("topic.1");
            topic.receive(message);
        } catch (Exception e) {
            System.err.println("caught an exception: " + e);
            e.printStackTrace();
        }
    }
}
```

Figure 10.8: This program uses the registry to locate the remote object that is named `topic.1` and that implements the `Topic` interface. The program then asks that object to receive a message.

method, and deletions happen when a message cannot be successfully delivered. Because the RMI infrastructure is allowed to invoke each operation in its own thread, the remote operations are marked as **synchronized** so as to provide mutual exclusion. This prevents any races in the manipulations of the list of subscribers. When the **TopicServer** program is run from the command line, the **main** method creates an instance of the class, exports it for remote access, and places it in the local registry, using the same **topic.1** name as the **Publisher** class looks up.

The final component of the example publish/subscribe system is the **Subscriber** class, as shown in Figure 10.11. This class provides a simple test program which displays all the messages it receives. Like the **Publisher** class, it uses the registry on a specified host or on the local host if none is specified. Also like the **Publisher** class, it looks up the name **topic.1** in that registry, thereby obtaining a reference to some remote object implementing the **Topic** interface. The reference will actually be to a proxy that implements the interface. However, the proxy will be communicating with an instance of the **TopicServer** class. Unlike the **Publisher**, the **Subscriber** is itself a remotely accessible object. It is created and exported just like the **TopicServer** is. However, it is not bound in the registry; the **TopicServer** does not locate its subscribers by name.

Before you can successfully run the **TopicServer** and test it using the **Publisher** and **Subscriber** programs, you will probably need to run the **rmiregistry** program that comes as part of the Java system. The details of how you run this program are system-specific, as is the mechanism for ensuring that all components of the overall RMI system have access to your classes. Therefore, you are likely to need to consult the documentation for your specific Java system in order to successfully test the sample code or complete the programming projects. Once you get over these technical hurdles, however, you will be able to communicate among multiple machines, so long as they are all running Java and so long as no network firewalls impede communication among them. In the following section, you will see how web services provide an alternate RPC mechanism that can allow communication between an even wider assortment of machines.

10.4 Web Services

A *web service* is a communicating component that complies with a collection of Internet standards designed to share as much as possible with the standards used for ordinary web browsing. This allows web services to take

```
import java.rmi.registry.Registry;
import java.rmi.registry.LocateRegistry;
import java.rmi.RemoteException;
import java.rmi.server.UnicastRemoteObject;
import java.util.List;
import java.util.ArrayList;

public class TopicServer implements Topic {

    private List<MessageRecipient> subscribers;

    public TopicServer(){
        subscribers = new ArrayList<MessageRecipient>();
    }

    public synchronized void receive(String message)
        throws RemoteException
    {
        List<MessageRecipient> successes =
            new ArrayList<MessageRecipient>();
        for(MessageRecipient subscriber : subscribers) {
            try {
                subscriber.receive(message);
                successes.add(subscriber);
            } catch(Exception e) {
                // silently drop any subscriber that fails
            }
        }
        subscribers = successes;
    }
}
```

Figure 10.9: The `TopicServer` class continues in the Figure 10.10. The `receive` method shown here is remotely invoked by publishers and itself remotely invokes the `receive` method of subscribers.


```
public synchronized void addSubscriber(MessageRecipient subscriber)
    throws RemoteException
{
    subscribers.add(subscriber);
}

public static void main(String args[]) {
    try {
        TopicServer obj = new TopicServer();
        Topic stub =
            (Topic) UnicastRemoteObject.exportObject(obj, 0);
        Registry registry = LocateRegistry.getRegistry();
        registry.rebind("topic.1", stub);
        System.err.println("Server ready");
    } catch (Exception e) {
        System.err.println("Server exception: " + e.toString());
        e.printStackTrace();
    }
}
```

Figure 10.10: This continuation of the `TopicServer` class, begun in Figure 10.9, shows how remote objects are created, exported (that is, made remotely accessible), and bound to a name in the registry.

```
import java.rmi.registry.LocateRegistry;
import java.rmi.registry.Registry;
import java.rmi.RemoteException;
import java.rmi.server.UnicastRemoteObject;

public class Subscriber implements MessageRecipient {

    public synchronized void receive(String message)
        throws RemoteException
    {
        System.out.println(message);
    }

    public static void main(String[] args) {
        if(args.length > 1){
            System.err.println("Usage: java Subscriber [hostname]");
            System.exit(1);
        }
        String hostname = args.length > 0 ? args[0] : null;
        try {
            Registry registry = LocateRegistry.getRegistry(hostname);
            Topic topic = (Topic) registry.lookup("topic.1");
            Subscriber obj = new Subscriber();
            MessageRecipient stub = (MessageRecipient)
                UnicastRemoteObject.exportObject(obj, 0);
            topic.addSubscriber(stub);
        } catch (Exception e) {
            System.err.println("caught an exception: " + e);
            e.printStackTrace();
        }
    }
}
```

Figure 10.11: Instances of the `Subscriber` class are created and exported the same way as `TopicServers` are, so that they can be remotely accessed. However, they are not bound in the registry. Instead, the stub referring to the `Subscriber` is passed to the `addSubscriber` method so the `TopicServer` can store the reference away for later.

advantage of the web's popularity, hopefully making communication between programmed components as ubiquitous as the communication with humans facilitated by the web.

The web services standards are based on *XML (Extensible Markup Language)*, a form of structured textual document. All XML documents have nested components with explicitly indicated types and attributes. The specific kinds of nested components depend on the XML application. For example, where XML is used for request messages, the component parts indicate what operation is being invoked and what arguments are being supplied to it. By contrast, where XML is used not to invoke an operation but to define an interface, the component parts enumerate what the interface's operations are, what kinds of messages should be exchanged to invoke those operations, and so forth.

Web service interfaces are described using an XML notation known as *WSDL (Web Services Description Language)*. This notation is rather verbose and is not usually read or written by humans. Instead, the humans normally use user-friendly tools to process the WSDL, which serves as a common interchange format accepted by all the tools. However, you can get a feel for WSDL by looking at the excerpts shown in Figure 10.12. The GoogleSearch API, from which these are taken, provides operations for searching for web pages. However, it also provides an operation for suggesting spelling corrections, as shown here. The operation involves two message types, one for request messages and one for response messages. Request messages contain two string arguments; one is an access control key that Google demands so as to limit use of their service, and the other is the phrase to correct. The response message contains just the returned value, a string containing the suggested correction.

Notice that in Figure 10.12, the `doSpellingSuggestion` operation is explicitly specified as using an input request message and an output response message. Because WSDL provides this detailed specification of how operations exchange messages, it can be used for patterns of communication other than RPC. The most common usage is for RPC, as with the GoogleSearch API. However, an operation can have only an input message, in which case the web service fits the messaging model instead of the RPC model. In theory, an operation could also specify a more extended message exchange pattern, with multiple input and output messages; however, I am unaware of any use of this feature.

The WSDL standard allows providers of web services to make their interfaces known to potential users without concern for what programming language or implementation technology they use. For example, I cannot tell

```
<message name="doSpellingSuggestion">
  <part name="key"           type="xsd:string"/>
  <part name="phrase"        type="xsd:string"/>
</message>

<message name="doSpellingSuggestionResponse">
  <part name="return"        type="xsd:string"/>
</message>

<operation name="doSpellingSuggestion">
  <input message="typens:doSpellingSuggestion"/>
  <output message="typens:doSpellingSuggestionResponse"/>
</operation>
```

Figure 10.12: These excerpts from the WSDL definition of the GoogleSearch API show the two messages used to ask for and receive a spelling suggestion and the operation that combines those two messages.

from the WSDL excerpted in Figure 10.12 whether Google is using J2EE, Microsoft's .NET, or some other technology. I am free to use whichever I choose in writing my client.

For this goal of interoperability to be realized, the service providers and users need to agree on more than just WSDL as a means of specifying interfaces. They also need to agree upon the specific format for transmitting the request and response messages. For this purpose, web services use a second XML format, known as *SOAP*. (SOAP once stood for Simple Object Access Protocol but no longer does.) Each SOAP document is a message and should match one of the message descriptions from the WSDL interface description. For example, you saw WSDL message descriptions for the two message types `doSpellingSuggestion` and `doSpellingSuggestionResponse`. Figures 10.13 and 10.14 show specific SOAP messages that fit these two descriptions. The first one is a message asking for suggestions as to how “middlewear” should really be spelled, and the second is a message responding with the suggestion of “middleware.”

Some transport mechanism needs to underlie SOAP. That mechanism delivers the string of bytes shown in Figure 10.13 to the server and then deliver the bytes shown in Figure 10.14 back to the client. The most common transport mechanism for SOAP is HTTP, the application-layer protocol normally used to access web pages. Notice that in web services terminology, HTTP is

```
<?xml version="1.0" encoding="UTF-8"?>
<env:Envelope
  xmlns:env="http://schemas.xmlsoap.org/soap/envelope/"
  xmlns:xsd="http://www.w3.org/2001/XMLSchema"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:enc="http://schemas.xmlsoap.org/soap/encoding/"
  xmlns:ns0="urn:GoogleSearch"
  env:encodingStyle="http://schemas.xmlsoap.org/soap/encoding/">
  <env:Body>
    <ns0:doSpellingSuggestion>
      <key xsi:type="xsd:string">GoogleAccessControlKeyHere</key>
      <phrase xsi:type="xsd:string">middlewear</phrase>
    </ns0:doSpellingSuggestion>
  </env:Body>
</env:Envelope>
```

Figure 10.13: This example SOAP message asks Google for spelling suggestions on the string `middlewear`. This message has been broken into indented lines for legibility and has a place-holder where a real message would contain an access-control key issued by Google, which I am not allowed to divulge.

```
<?xml version='1.0' encoding='UTF-8'?>
<SOAP-ENV:Envelope
  xmlns:SOAP-ENV="http://schemas.xmlsoap.org/soap/envelope/"
  xmlns:xsi="http://www.w3.org/1999/XMLSchema-instance"
  xmlns:xsd="http://www.w3.org/1999/XMLSchema">
  <SOAP-ENV:Body>
    <ns1:doSpellingSuggestionResponse
      xmlns:ns1="urn:GoogleSearch"
      SOAP-ENV:encodingStyle=
        "http://schemas.xmlsoap.org/soap/encoding/">
      <return xsi:type="xsd:string">middleware</return>
    </ns1:doSpellingSuggestionResponse>
  </SOAP-ENV:Body>
</SOAP-ENV:Envelope>
```

Figure 10.14: This example SOAP message returns `middleware` as a spelling suggestion in response to `middlewear`. (Line breaks and indentation changed for legibility.)

referred to as a transport, because it conveys the SOAP messages, whereas in traditional networking terminology, the transport layer is one layer lower, where TCP operates. In effect, web services are building a super-application-layer on top of the application layer, thereby treating the HTTP application layer as though it were only a transport layer. As mentioned in Chapter 9, one advantage of this arrangement is that it circumvents obstacles such as firewalls that stand in the way of deploying new application-layer protocols. Almost any Internet connection is open to HTTP traffic.

When HTTP is used for a request-response message pair, as in the spelling suggestion example, the client opens a connection to the server exactly as it would to an ordinary web server, providing a URL that represents the particular web service, known as an *endpoint address*. The client then sends the SOAP request message as the body of a POST method, the kind of HTTP transaction more traditionally used for filled-in forms. The server sends the SOAP response message in the body of its HTTP response.

Although SOAP is most commonly used with HTTP, the web services architecture is intentionally neutral with regard to transport. SOAP messages can equally well be sent as the bodies of email messages, using SMTP, or as messages in a reliable message-queuing system, such as WebSphere MQ.

If you remember that the goal of communication middleware is to ease application programmers' burdens, it should be obvious that SOAP and WSDL are not intended to be used without the aid of automated tools. You could, in principle, read the GoogleSearch API's WSDL specification yourself and based on it write code that sent the SOAP message shown in 10.13 over HTTP. You could do this by using nothing more than the ability to open up a TCP socket and send bytes through it. Then you could read in from the socket the bytes constituting Figure 10.14's response and arduously extract from it the spelling suggestion being returned. However, this would be making a distributed system harder to construct, not easier.

Luckily, there are a variety of language-specific and vendor-specific tools that make web services much easier to construct. In particular, both .NET and J2EE have support for web services. As an example, let's look at *JAX-RPC* (*Java API for XML-Based RPC*), a component of J2EE.

In Programming Project 10.10, you can use a JAX-RPC tool to automatically translate the GoogleSearch WSDL file into a Java interface that contains ordinary Java-callable methods for each of the web service's operations. For example, it contains

```
public String doSpellingSuggestion(String key, String phrase);
```

Using this, you can set a variable to the suggested spelling with just this code:

```
suggestion = aGoogleSearch.doSpellingSuggestion(  
    "GoogleAccessControlKeyHere",  
    "middleware");
```

The Java object named `aGoogleSearch` is a stub proxy implementing the interface created from the WSDL file; a few prior lines of code would set it up. This proxy takes care of generating the big, messy SOAP request message, sending it, reading in the response, and picking the suggestion out from amid all its SOAP wrappings. You, the application programmer, don't need to do any of that.

The WSDL and SOAP facilities described thus far provide the core facilities for web services, but there are many other standards, and proposed standards, for less central aspects. The entire area is in considerable flux with many competing draft standards. However, one other standard is approximately as solid as WSDL and SOAP are. That standard, *UDDI* (*Universal Description, Discovery, and Integration*), provides a way for web service providers to list their services in a registry and for potential users to discover them by looking in the registry for services matching a description. UDDI registries are themselves web services, accessed via SOAP messages in accordance with a WSDL specification.

10.5 Security and Communication Middleware

Messaging systems and RPC servers often use ACLs to control access, much like file systems do. For example, a broker with a hierarchy of publish/subscribe topics can associate an ACL with each topic in the hierarchy; the ACL specifies the users or groups that may publish and those that may subscribe. ACLs on subtopics take precedence over those on more general topics. Thus, protection can be specified as precisely as necessary for those subtopics where it matters while allowing most subtopics the convenience of inheriting an ancestor topic's ACL.

An ACL lists the users or groups that should be granted access, but this still leaves open one of the most difficult aspects of security in a distributed system. Namely, how should a server know which user's access rights apply for each incoming connection? Any robust solution to this problem relies on the cryptographic mechanisms described in Section 9.6. I can illustrate this using an example from web services.

Recall that the exchange of SOAP messages between a client and web service normally takes place using the same HTTP protocol as is used for browsing the web. As such, the same cryptographic security mechanisms are used by interposing the Secure Sockets Layer, SSL, between HTTP and the underlying TCP connection.

Just as with a secure web site, a secure web service identifies itself by using a *certificate*, which is a document attesting to the server's identity and to the public half of the server's asymmetric key pair. This public key can be used by the client to check the server's digitally signed messages and also to send the server a secret key for confidential communication. The certificate itself is digitally signed by some trusted *Certification Authority* (CA), an organization that has made its public key well known and that can be counted on to check out the legitimacy of another organization's or individual's identity claim before issuing a certificate.

The server's certificate allows the client to trust that it is communicating with the real server and not an impostor. However, the server still has no idea which user identity to associate with the client. Two options exist for solving that problem, one that continues to follow the lead of ordinary web sites used by humans and another that may be better suited to widespread deployment of web services. I will present the solution first that you are probably familiar with from your own use of the web and then the more robust alternative.

When you connect to a secure web site, your browser checks the server's certificate and if all is well signals this fact by showing you a locked padlock. The server then typically asks you to enter a username and password for the site. The strings that you enter are sent over the SSL-encrypted communication channel and so are not subject to eavesdropping or tampering in transit. Moreover, because your browser checked the server's certificate, you can be sure you aren't sending your password to a con artist. The server gets the strings in decrypted form and checks them against its user database. This style of authentication relies on you and the site having a shared secret, the password. In general, each client/server pair requires a shared secret established in advance.

This first style of client authentication, which is provided by HTTP under the name *basic authentication*, can be a workable method for web services that are not widely deployed, especially for those that are deployed only internally to an enterprise. In that context, the various web services will ordinarily be controlled by the same administrators and as such can all share a common authentication server that keeps track of users with their passwords. Thus, a secret password needs to be established for each user,

but not for each user/service pair. Even across enterprise boundaries, basic authentication may suffice for web services that have only a small number of users, such as a web service used to facilitate one particular relationship between a pair of enterprises.

Before I move on to the more sophisticated alternative, it is worth contrasting the first alternative, basic authentication using SSL, with weaker password-based authentication. Consider, for example, the GoogleSearch API's spelling suggestion operation, which was shown in Section 10.4. This operation takes a secret access-control key as an argument in the request message itself. The access-control key is issued by Google and essentially acts as a combination of username and password in a single string. However, the GoogleSearch web service does not use SSL; it uses ordinary unencrypted HTTP directly over TCP. One consequence is that the access control keys are subject to eavesdropping and so could be captured and then reused. However, there is a second way in which a malefactor could capture a key.

Recall that with SSL, the client program receives a certificate of the server's identity, protecting it against impostors. Because GoogleSearch is not using SSL, you could be sending your misspellings to an impostor, perhaps someone who wants to embarrass you. Moreover, because you send your key along, you could also be sending your key to an impostor. This helps explain the significance of SSL's server authentication. It not only protects the client from rogue servers, but also protects the server from misuse through password capture. Even if you don't care whether your misspellings become public knowledge, Google presumably cares that their service isn't used indiscriminately. Otherwise they would not have established access-control keys.

What can you conclude, then, about Google's security design? Presumably they decided that their service was valuable enough to make some attempt to discourage casual misuse, but not so valuable that they were willing to pay the price of SSL cryptography to keep determined adversaries away. Also, their main concern is not with the actual identity of the user, but with limiting the number of searches made by any one user. If someone captures your GoogleSearch key, they will simply share your daily limit on searches, not imposing any extra burden on Google. Thus, Google's design stems from a well thought-out cost-benefit analysis, paradigmatic of how security decisions ought to be made. They did not make a mistake in using passwords without SSL. However, you would be making a mistake to blindly emulate their security mechanism on a web service of greater value.

Let us return, then, to security mechanisms suitable for high-value targets that are likely to attract serious attackers. Recall that the problem with

using HTTP's basic authentication over SSL is that it requires a shared secret password for each pair of client and server. If web services are to form a ubiquitous Internet-wide economy, as some prognosticators suggest, this will not be workable. Any client must be able to securely access any web service without prearrangement.

To solve this problem, web services can use the *mutual authentication* feature of SSL, which is almost never used for ordinary human-oriented web sites. In mutual authentication, both the client and the server have digitally-signed certificates obtained from trusted Certification Authorities. They exchange these certificates in the initial setup handshake that starts an SSL connection. Thus, without needing any usernames or passwords, each knows the identity of its communication partner for the entire duration of the connection. Mutual authentication is impractical for ordinary consumer-oriented web browsing, because merchants don't want to require all their customers to go to the trouble and expense of getting certificates. However, for the business-to-business communications where web services are expected to play their major role, mutual authentication seems well suited. It does, however, still have limitations, as I explain next.

The use of SSL, sometimes with mutual authentication, is widely deployed in practical web service applications. It also is incorporated in the most mature standards for web services; the Web Services Interoperability Organization's Basic Profile specifies that web service instances may require the use of HTTP over SSL and, in particular, may require mutual authentication. However, this approach to web service security has a fundamental limitation, with the result that more sophisticated, less mature standards take a different approach. The fundamental limitation is that SSL secures communication channels, rather than securing the SOAP messages sent across those channels.

To understand the difference between securing a channel and securing a message, consider the fact that several SOAP messages, originated by different users running different applications, may be sent across a single network connection. Consider also the fact that a single SOAP message may be relayed through a whole succession of network connections on its way from its originator to its ultimate destination. In both cases, SSL allows each computer to be sure of the authenticity of the neighboring computer, but it doesn't provide any direct support for associating the SOAP message with its author.

Therefore, the Web Services Security standard provides a mechanism whereby the XML format of a SOAP message can directly contain a digital signature for the message. Standards also govern the transmission of XML

in encrypted form and the use of XML to send certificates. Using these mechanisms, the Web Services Interoperability Organization's Basic Security Profile, currently only a working draft, provides requirements for SOAP messages to be encrypted and digitally signed. Because the signature is on the message itself, it can be forwarded along with the message through relay channels and has sufficient specificity to allow messages of varying origins to share a network connection.

One final advantage of the Web Services Security approach compared with SSL is that SOAP messages that have been digitally signed support non-repudiation, as described in Section 9.6. That is, because the recipient is in no better position to forge a message than anyone else would be, the recipient can present the message to a third party with a convincing argument that it came from the apparent sender. Today, web services are largely used within organizations and between close business partners with a high degree of mutual trust. However, as web services spread into more arms-length dealings between parties that have no established relationship, non-repudiation will become more important. Moreover, even if the communicating enterprises have a trust relationship, individual employees may be corrupt; digital signatures limit the scope of investigation that is needed if tampering is suspected.

Exercises

- 10.1 How does messaging differ from sending bytes over a TCP connection?
- 10.2 How does messaging differ from sending an email message?
- 10.3 How does messaging differ from RPC?
- 10.4 Does using response messages turn a message-queuing system into the equivalent of an RPC system? Why or why not?
- 10.5 Are web services an alternative to messaging and RPC systems, that is, a third kind of communication middleware? Why or why not?
- 10.6 For each of the following communication methods, give one example application scenario where you think it would be appropriate: message queuing, publish/subscribe messaging, RPC. In each case, justify your choice of communication method.
- 10.7 Recall that in publish/subscribe topic hierarchies, the wildcard `+` represents one component topic, whereas `#` represents a sequence of zero or

more components separated by slashes. Suppose a publish/subscribe system has topics **a**, **b**, **a/c**, **a/d**, **b/c**, **b/e**, **a/c/e**, and **a/d/e**. For each of the following subscriptions, specify which of those topics would be included: **a**, **a/+**, **a/#**, **a/c/+**, **a/+/e**, **#/e**.

- 10.8 Suppose **s** is a JMS messaging session and **d** is a JMS messaging destination. Show how to create a **Consumer** that would receive all messages sent to **d** containing a **Symbol** of **IBM** and that would also receive all those containing a **Price** of 0, independent of their **Symbol**.
- 10.9 In the RMI programming example, suppose several **Subscriber** objects are all subscribed to a single **TopicServer** and that several **Publisher** objects send messages to that **TopicServer**. Will all the **Subscribers** necessarily print the messages in the same order? Explain why or why not.
- 10.10 In the **TopicServer** implementation shown in Figures 10.9 and 10.10 on pages 456 and 457, the **receive** method invokes each subscriber's **receive** method. This means the **TopicServer**'s **receive** method will not return to its caller until after all of the subscribers have received the message. Consider an alternative version of the **TopicServer**, in which the **receive** method simply places the message into a temporary holding area and hence can quickly return to its caller. Meanwhile, a separate thread running in the **TopicServer** repeatedly loops, retrieving messages from the holding area and sending each in turn to the subscribers. What Java class from Chapter 4 would be appropriate to use for the holding area? Describe the pattern of synchronization provided by that class in terms that are specific to this particular application.
- 10.11 The text shown in Figure 10.15 has the right form to be a legal SOAP message, but it would not be legitimate to send this message to the **GoogleSearch** web service. Why not?
- 10.12 Section 10.5 mentions one reason why mutual authentication using certificates is not common in the human-oriented web: merchants don't want to turn customers off by requiring them to get certificates. One item of context here is that most consumers do business with only a small number of merchants. This is starting to change, as more businesses develop online presences and as consumers start branching out and shopping online for more than just books, music, videos, and airline tickets. Can you see any reason why this change might

```

<?xml version="1.0" encoding="UTF-8"?>
<env:Envelope
  xmlns:env="http://schemas.xmlsoap.org/soap/envelope/"
  xmlns:xsd="http://www.w3.org/2001/XMLSchema"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:enc="http://schemas.xmlsoap.org/soap/encoding/"
  xmlns:ns0="urn:GoogleSearch"
  env:encodingStyle="http://schemas.xmlsoap.org/soap/encoding/">
  <env:Body>
    <ns0:doSpellingSuggestion>
      <key xsi:type="xsd:int">314159</key>
      <phrase xsi:type="xsd:string">middlewear</phrase>
    </ns0:doSpellingSuggestion>
  </env:Body>
</env:Envelope>

```

Figure 10.15: This is a legal SOAP message but is not legitimate for sending to the GoogleSearch web service.

affect consumers' willingness to acquire certificates rather than use passwords?

Programming Projects

- 10.1 Create an RMI analog of the message-storage server of Figure 9.7 on page 411 and its companion client of Figure 9.8 on page 413.
- 10.2 Modify the `TopicServer` class shown in Figures 10.9 and 10.10 on pages 456 and 457 as described in Exercise 10.10. Be sure to correctly synchronize access to the list of subscribers.
- 10.3 Exercise 10.10 describes one way to modify the `TopicServer` class so that the `receive` method does not need wait for each subscriber's `receive` method, at least under normal circumstances. An alternative design to achieve that same goal would be for the `TopicServer`'s `receive` method to create a new thread for each incoming message. The thread would deliver that one message to the subscribers. Modify the `TopicServer` class shown in Figures 10.9 and 10.10 on pages 456 and 457 in this alternate way. Be sure to correctly synchronize access

to the list of subscribers.

- 10.4 In the RMI example code given in Section 10.3.2, only a single topic is used, bound in the registry to the name `topic.1`. Show how the `Publisher`, `TopicServer`, and `Subscriber` programs can be generalized to take a topic name as an additional command line argument, with each topic separately bound in the registry. Demonstrate the concurrent execution of two different topic objects on the same host, each with its own subscribers.
- 10.5 In Programming Project 10.4, you accommodated multiple publish/subscribe topics by having a separate `TopicServer` for each and by registering each in the registry. An alternative design would be to have a single `TopicServer` object, but with the `receive` and `addSubscriber` methods taking an extra argument that is the topic name. Develop and demonstrate the code for this approach. You may want to include extra methods for such purposes as adding topics and obtaining a list of the current topics.
- 10.6 The publish/subscribe system provided as an RMI example in Section 10.3.2 does not include a method for removing a subscriber from a topic. Arguably, such a method would be redundant, because the `TopicServer` class is prepared for subscribers that fail. A subscriber that wishes to unsubscribe could simply arrange to intentionally fail. However, that option doesn't handle the case of a subscriber that is subscribed to more than one topic and wishes to unsubscribe from one of them. The design would be cleaner and more flexible if the `Topic` interface and `TopicServer` class supported a `removeSubscriber` method. Add one and demonstrate its use.
- 10.7 Section 10.3.2 shows how RMI can be used to convey textual messages from publishers to subscribers by way of intermediate topic objects. If you have the requisite skill in building user interfaces in Java, you could use this RMI mechanism as the foundation for a chat-room application or a multi-player game. Write such a program. Depending on your design, you may want to incorporate some of the features from earlier programming projects; for example, multiple topics could support multiple chat rooms. You are also welcome to change the message type; an application-specific class of game moves might be more appropriate than textual strings.

- 10.8 The `Publisher` class in Figure 10.8 on page 454 makes use of the `Topic` interface even though the `MessageRecipient` interface would suffice. Change the class to use the more general interface and demonstrate that, with appropriate changes elsewhere, the `Publisher` can wind up communicating either directly with a `Subscriber` or with an intermediary `TopicServer` as before.
- 10.9 The `Topic` interface in Figure 10.7 on page 453 extends the base interface `MessageRecipient` and also uses that same interface as the argument type for the `addSubscriber` method. Demonstrate how this allows one `TopicServer` to function as a subscriber to another `TopicServer`. Assuming that you have done Programming Project 10.4, there is no need for the `TopicServer` that is functioning as a subscriber to add itself to the other one using `addSubscriber`. Instead, you can leave the code for `TopicServer` unchanged and add a new program that looks up the two `TopicServer` objects in the registry and adds one as a subscriber to the other.
- 10.10 Acquire an access control key for `GoogleSearch` from Google and download the software associated with the *J2EE 1.4 Tutorial*. After working through the JAX-RPC portion of the tutorial, modify one of the example clients so that it gets a spelling suggestion from `GoogleSearch` instead of accessing the example `Hello` web service. You can use `http://api.google.com/search/beta2` as the endpoint address and `http://api.google.com/GoogleSearch.wsdl` as the WSDL location. Optionally, you can use a packet capture program such as `wireshark` to verify that the web service is being accessed through ordinary HTTP, without the use of SSL, and that the SOAP messages are essentially as shown in Figures 10.13 and 10.14.

Exploration Projects

- 10.1 Read about message-driven beans in the *J2EE 1.4 Tutorial* and write a concise explanation of what they are and why they are more convenient than directly using JMS.
- 10.2 Work through the examples in Chapters 28 and 33 of the *J2EE 1.4 Tutorial*, “A Message-Driven Bean Example” and “The Java Message Service API.”

Notes

The topics in this chapter are subject to particularly rapid technical developments. As such, your best source of information is likely to be the web sites. The Java web site, <http://java.sun.com>, has information both on RMI and on J2EE, which includes JMS and JAX-RPC. The Web Services Activity web site, <http://w3c.org/2002/ws/>, has information on WSDL, SOAP, and web services in general. Other important sites for web services standards are the Web Services Interoperability Organization, <http://www.ws-i.org/>, and OASIS, <http://www.oasis-open.org/>, which tends to have more specialized, advanced standards. The information on these sites—and in many published books for that matter—tends to emphasize the technical details over the big picture of how to use the technology. One book that does provide a lot of big-picture advice on the use of messaging is by Hohpe and Woolf [78].